

# EMBEDDINGS 2



# WHAT IS FASTTEXT?

FastText treats each word as composition of n-grams:

example: ball

"<ba"

"bal"

"ball"

"all"

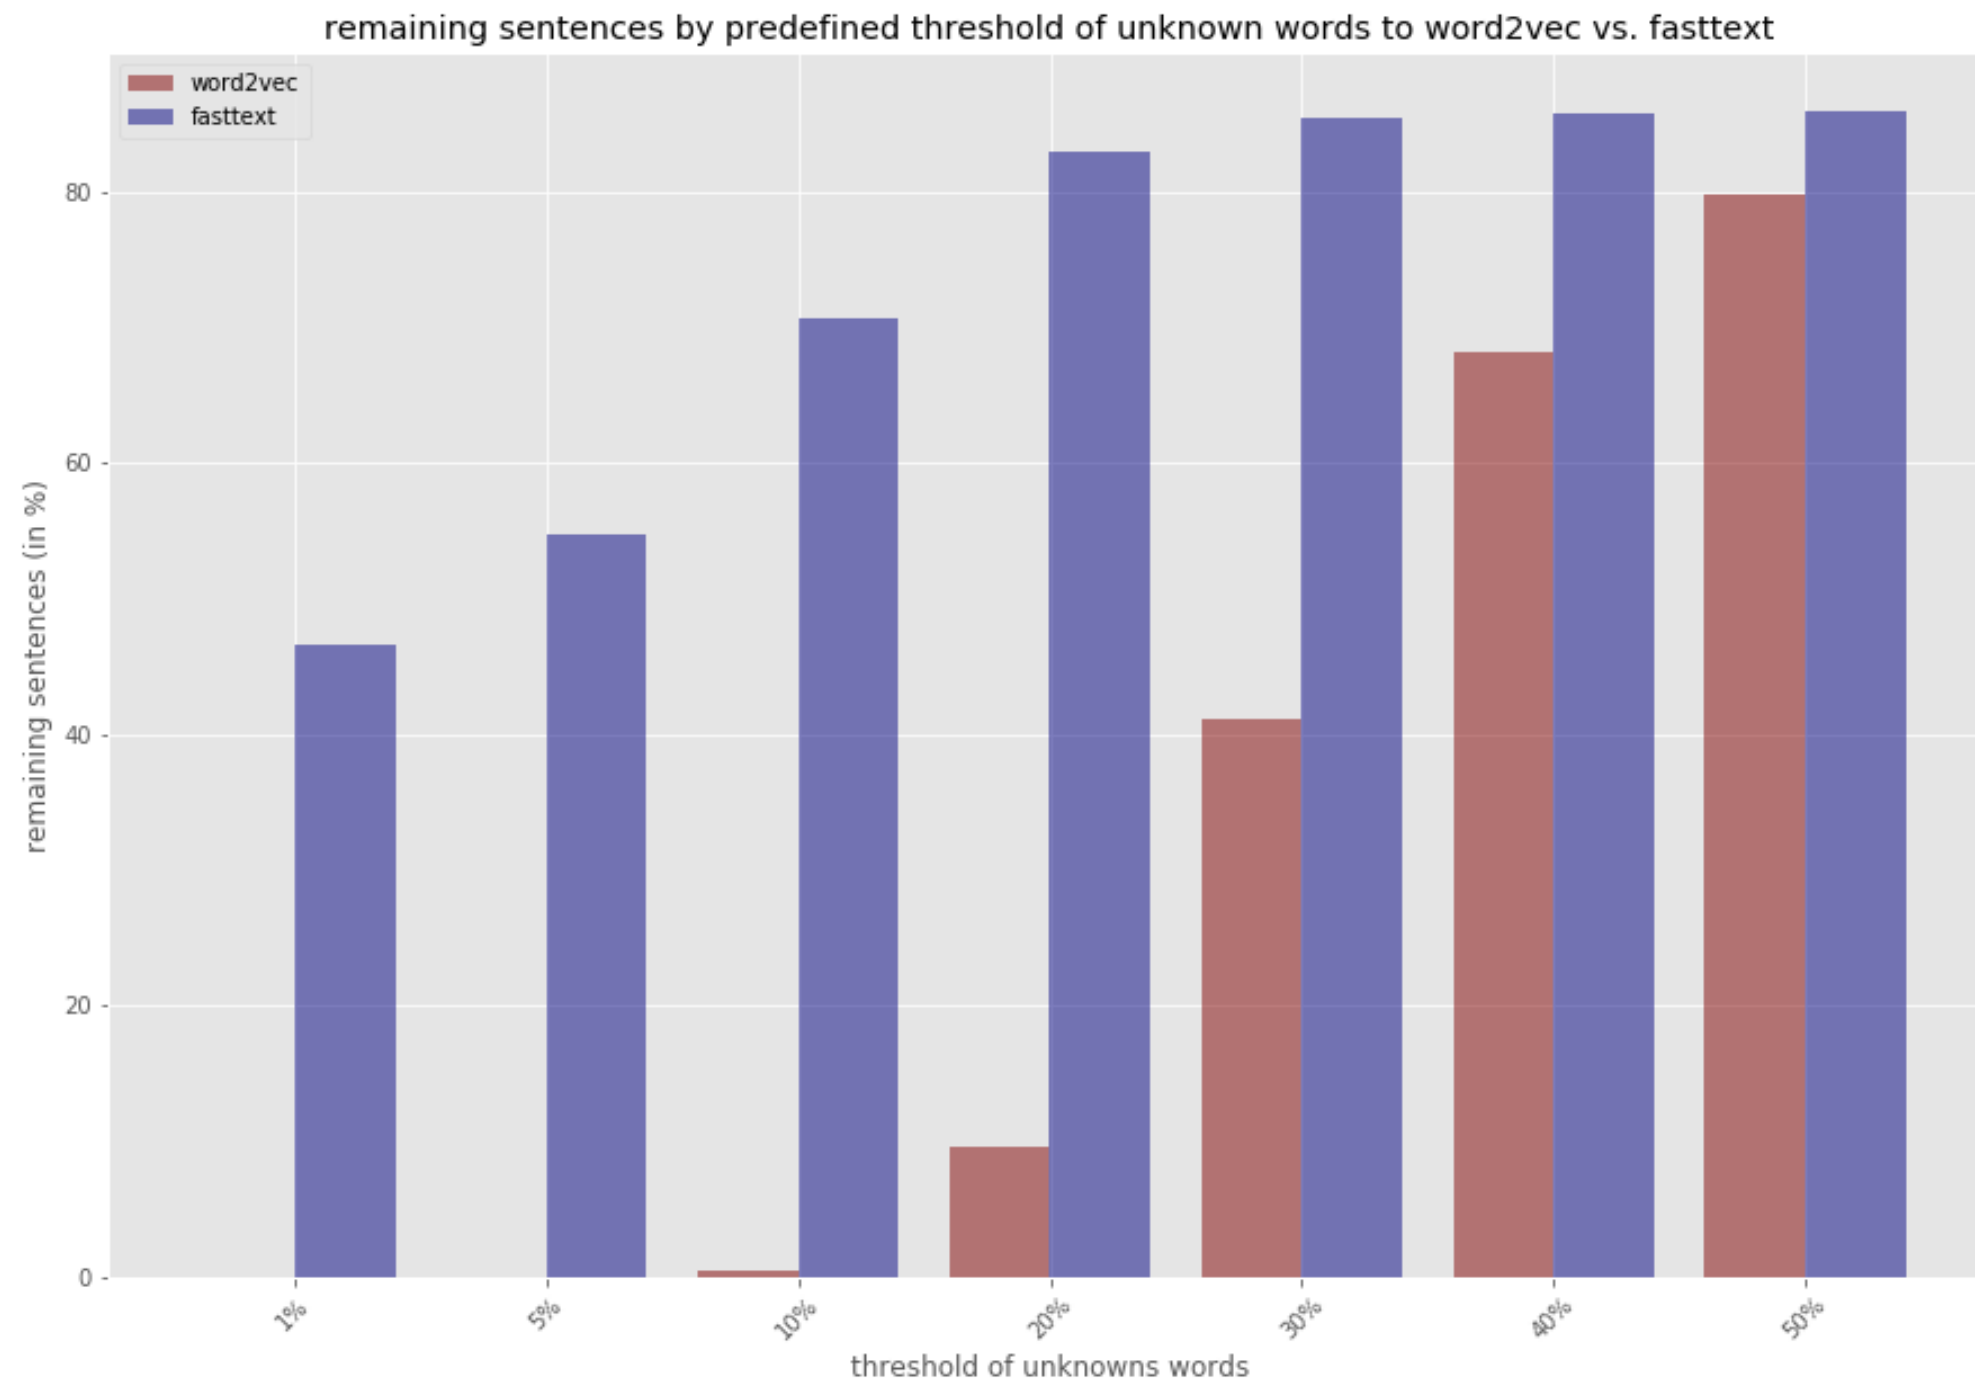
"||>"

with smallest n-grams set to 3 and highest to 4

Note that the sequence <all>, corresponding to the word *all* is different from the tri-gram all from the word *ball*.



# WORD2VEC VS. FASTTEXT





# FASTTEXT.ZIP

## COMPRESSING TEXT CLASSIFICATION MODELS

What?

- reduce memory of text classification model

How?

- quantization
- retraining
- vocabulary pruning
- hashing



# QUANTIZATION

The purpose of vector quantization is to compress vectorial data.

general idea:

- find a good set of reference vectors
- replace each data vector by the **index** of its best reference vector

# RETRAINING

- retrain the layers occurring after the quantization,  
so that the network can re-adjust itself to the quantization



# MEMORY SAVINGS

- compression factor of 10

(without any noticeable loss of performance)

- **without RETRAINING:**  
drop in accuracy of 0.5%



# VOCABULARY PRUNING

- feature selection problem:  
selecting  $K$  words & ngrams that preserve the highest magnitude to the model
- BUT: some documents won't have any of the  $K$  best features

 greedy approach

each document -> already covered by retained feature?  
if not: add highest magnitude feature to set of retained features



# HASHING


- *John likes to watch movies.*
- *Mary likes movies too.*
- *John also likes football.*

Term	Index
John	1
likes	2
to	3
watch	4
movies	5
Mary	6
too	7
also	8
football	9

John	likes	to	watch	movies	Mary	too	also	football
1	1	1	1	1	0	0	0	0
0	1	0	0	1	1	1	0	0
1	1	0	0	0	0	0	1	1



# MEMORY SAVINGS

- compression factor of up to 100  
(dependent on the dataset used)
  - complementary with quantization!
-  combination of both strategies:  
model size reduction of up to x1000



# ELMO

What?

- word reps that model:
  - (1) complex characteristics of word use  
e.g. syntax, semantics
  - (2) how this word use vary across context

How?

- word rep as function of entire input sentence
- computed on biLMs with character convolution
- semi-supervised learning



► ELMo: Embeddings from Language Models



# USE OF ENTIRE INPUT SENTENCE

- each token is assigned a representation

that is a function

of the entire input sentence



sentence level embeddings surpass the performance  
of using word level embeddings



# BILM

- vectors are derived from a bidirectional LSTM that is trained with a coupled language Model (LM)

The diagram shows the sentence "Rubin, how are you?". The word "are" is the target word, indicated by a blue vertical line and the label "Target Word" above it. A red bracket under "Rubin, how" is labeled "Left Context". A green bracket under "you?" is labeled "Right Context".

Predicting the word "are" from both left and right contexts.



# ELMO REPRESENTATIONS ARE ?DEEP?

- they are a function of all internal layers of a biLM
- learning a linear combination of the vectors stacked behind each input word (improves over just using the top LSTM layer)
- higher-level LSTM states capture context-dependent aspects of word meaning (semantics)
- lower-level LSTM states model aspects of syntax
- exposing these signals allows semi-supervision



# FUNCTIONALITY OF ELMo

- ELMo is a combination of the intermediate layer representations in the biLM
- For each token  $\mathbf{t}$ , a biLM with  $L$  layers computes a set of  $2L+1$  representations
- ELMo collapses all layers of the biLM into a single vector

$$\mathbf{ELMo}_k^{task} = \gamma^{task} \sum_{j=0}^L s_j^{task} \mathbf{h}_{k,j}^{LM}.$$

each LSTM layer outputs a context-dependent representation:

$$\mathbf{h}_{k,j}^{LM} = [\vec{\mathbf{h}}_{k,j}^{LM}; \overleftarrow{\mathbf{h}}_{k,j}^{LM}]$$

where  $j = 1, \dots, L$



# BENEFITS OF ADDING ELMO

- 20% relative error reduction
- increase the sample efficiency:
  - less parameter updates
  - smaller training set size

e.g. a Semantic role labeling task

without ELMo:

- model reach maximum F1 at **486** epochs of training

with ELMo:

- maximum F1 at **10** epochs of training

- 1% of the training set has the same F1 as the baseline model with **10%** of the training set



# UNIVERSAL SENTENCE ENCODER

What?

- encoding sentences into embedding vectors  
-> facing NLP transfer tasks with sentence embeddings

How?

- two variants of encoding models  
(trade-off between accuracy & computational effort)
  - transformer architecture
  - deep averaging network (DAN)



# TRANSFORMER

TARGETING HIGH ACCURACY!

- designed for a general purpose

 multi-task learning:

- a Skip-Thought like task
- an input-response task
- a classification task (to train on supervised data)

Skip-Thought replaces LSTM!



# DEEP AVERAGING NETWORK (DAN)

TARGETING EFFICIENT COMPUTE RESOURCES!

- Averaging of word embeddings and bi-gram embeddings
- afterwards passed through a feedforward deep neural network

► multi-task learning again:

one single DAN encoder is used for all 3 tasks!

- Advantage to Transformer:  
compute time is linear in the length of the input sentence!



# IS TECHNOLOGY BIASED?

Target words	Attrib. words	Ref	GloVe		Uni. Enc. (DAN)	
			d	p	d	p
Eur.-American vs Afr.-American names	Pleasant vs. Unpleasant 1	<i>a</i>	1.41	$10^{-8}$	0.361	0.035
Eur.-American vs. Afr.-American names	Pleasant vs. Unpleasant from (a)	<i>b</i>	1.50	$10^{-4}$	-0.372	0.87
Eur.-American vs. Afr.-American names	Pleasant vs. Unpleasant from (c)	<i>b</i>	1.28	$10^{-3}$	0.721	0.015
Male vs. female names	Career vs family	<i>c</i>	1.81	$10^{-3}$	0.0248	0.48
Math vs. arts	Male vs. female terms	<i>c</i>	1.06	0.018	0.588	0.12
Science vs. arts	Male vs female terms	<i>d</i>	1.24	$10^{-2}$	0.236	0.32
Mental vs. physical disease	Temporary vs permanent	<i>e</i>	1.38	$10^{-2}$	1.60	0.0027
Young vs old peoples names	Pleasant vs unpleasant	<i>c</i>	1.21	$10^{-2}$	1.01	0.022
Flowers vs. insects	Pleasant vs. Unpleasant	<i>a</i>	1.50	$10^{-7}$	1.38	$10^{-7}$
Instruments vs. Weapons	Pleasant vs Unpleasant	<i>a</i>	1.53	$10^{-7}$	1.44	$10^{-7}$

Table 4: Word Embedding Association Tests (WEAT) for GloVe and the Universal Encoder. Effect size is reported as Cohen's d over the mean cosine similarity scores across grouped attribute words. Statistical significance is reported for 1 tailed p-scores. The letters in the *Ref* column indicates the source of the IAT word lists: (a) Greenwald et al. (1998) (b) Bertrand and Mullainathan (2004) (c) Nosek et al. (2002a) (d) Nosek et al. (2002b) (e) Monteith and Pettit (2011).



# RESOURCE USAGE

	Computation Time	Memory Usage
Transformer	$O(n^2)$	$O(n)$
DAN	$O(n^2)$	$O(1)$



# RESULTS

- In general: Transformer performs as good or better than DAN
- for some tasks: DAN performs as good or better than Transformer
- ▶ best performance on most tasks:  
models that make use of both sentence and word level transfer!
- transfer learning is most helpful if less training data is available
- for each NLP task the trade-off between accuracy & complexity (computation & memory) has to be considered!